

Week 2: EDA

Monica Alexander

19/01/2021

By the end of this lab you should know

- how to get the dimensions of a data frame (rows, columns)
- the `group_by` function and how to use it to get summary statistics by group
- how to filter out missing values (NA) in one column or multiple columns, using `!is.na()` or `drop_na()`
- how to calculate the correlation coefficient between two variables
- how to get the number of observations by group
- how to calculate proportions by group
- `ggplot` basics; how to make each of the important types of graphs
 - histogram
 - bar chart
 - boxplot
 - line plot
 - scatter plot
- how to color / fill by group
- `fct_reorder` to reorder categorical values
- selecting only certain values of a variable using `%in%`
- if there's time: faceting

Read in data

Generally there are 3 steps in setting up any R session: 1. Choose the packages you are going to use and tell R to equip them by the `library()` command. + We are using two additional packages today. You will see the use of `skimr` in a moment, but the `here` package allows us to access our files easier. Since we haven't use these packages before we need to install them. 2. Next we set our working directory. This tells R where all the files are and how to access them. Since we are using the `here` package as well, we are setting both our working directory and `here` package.

3. Read your files from the appropriate folder. Use `read_csv` command to read the file.

```
#install.packages("skimr") # Install new packages
#install.packages("here")

# Call the packages that you are using
library(tidyverse)
```

```
## -- Attaching packages -----
```

```
## v ggplot2 3.3.2      v purrr  0.3.4
## v tibble  3.0.3      v dplyr  1.0.2
## v tidyr   1.1.2      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.5.0
```

```
## -- Conflicts -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(skimr)
library(here)
```

```
## here() starts at /Users/monicaalexander/src/soc6707
```

```
# Read the file
```

```
gss <- read_csv(here("data/gss.csv")) # Only include data folder if your file is in a folder called data
```

```
## Parsed with column specification:
## cols(
##   .default = col_character(),
##   caseid = col_double(),
##   age = col_double(),
##   age_first_child = col_double(),
##   age_youngest_child_under_6 = col_double(),
##   total_children = col_double(),
##   age_start_relationship = col_double(),
##   age_at_first_marriage = col_double(),
##   age_at_first_birth = col_double(),
##   distance_between_houses = col_double(),
##   age_youngest_child_returned_work = col_double(),
##   feelings_life = col_double(),
##   hh_size = col_double(),
##   number_total_children_intention = col_double(),
##   number_marriages = col_double(),
##   fin_supp_child_supp = col_double(),
##   fin_supp_child_exp = col_double(),
##   fin_supp_lump = col_double(),
##   fin_supp_other = col_double(),
##   is_male = col_double(),
##   main_activity = col_logical()
##   # ... with 2 more columns
## )
```

```
## See spec(...) for full column specifications.
```

```
country_ind <- read_csv(here("data/country_indicators.csv"))
```

```
## Parsed with column specification:
## cols(
##   country_code = col_character(),
##   country = col_character(),
```

```
## region = col_character(),
## year = col_double(),
## tfr = col_double(),
## life_expectancy = col_double(),
## child_mort = col_double(),
## maternal_mort = col_double(),
## gdp = col_double()
## )
```

Overview summaries of data

Skim is useful for the GSS, because it gives a broad overview of what types of variables the dataset contains.

```
summary(country_ind)
```

```
## country_code      country      region      year
## Length:1584      Length:1584      Length:1584      Min.   :2009
## Class :character  Class :character  Class :character  1st Qu.:2011
## Mode  :character  Mode  :character  Mode  :character  Median :2013
##                                     Mean  :2013
##                                     3rd Qu.:2015
##                                     Max.   :2017
##      tfr      life_expectancy  child_mort      maternal_mort
## Min.   :1.131  Min.   :47.02  Min.   : 1.796  Min.   : 2.0
## 1st Qu.:1.748  1st Qu.:66.92  1st Qu.: 7.955  1st Qu.: 15.0
## Median :2.402  Median :75.78  Median :19.698  Median : 55.0
## Mean   :2.886  Mean   :73.25  Mean   :35.162  Mean   :179.7
## 3rd Qu.:3.984  3rd Qu.:79.70  3rd Qu.:55.609  3rd Qu.:248.0
## Max.   :7.511  Max.   :87.34  Max.   :166.192  Max.   :1450.0
##      gdp
## Min.   : 670.8
## 1st Qu.: 3591.3
## Median :10869.4
## Mean   :17371.6
## 3rd Qu.:24321.4
## Max.   :124024.6
```

```
skim(gss)
```

Table 1: Data summary

| | |
|------------------------|-------|
| Name | gss |
| Number of rows | 20602 |
| Number of columns | 85 |
| Column type frequency: | |
| character | 63 |
| logical | 1 |
| numeric | 21 |
| Group variables | None |

Variable type: character

| skim_variable | n_missing | complete_rate | min | max | empty | n_unique | whitespace |
|----------------------------|-----------|---------------|-----|-----|-------|----------|------------|
| sex | 0 | 1.00 | 4 | 6 | 0 | 2 | 0 |
| place_birth_canada | 97 | 1.00 | 10 | 19 | 0 | 3 | 0 |
| place_birth_father | 203 | 0.99 | 10 | 19 | 0 | 3 | 0 |
| place_birth_mother | 47 | 1.00 | 10 | 19 | 0 | 3 | 0 |
| place_birth_macro_region | 16457 | 0.20 | 4 | 18 | 0 | 6 | 0 |
| place_birth_province | 4289 | 0.79 | 6 | 39 | 0 | 11 | 0 |
| year_arrived_canada | 16550 | 0.20 | 19 | 27 | 0 | 14 | 0 |
| province | 0 | 1.00 | 6 | 25 | 0 | 10 | 0 |
| region | 0 | 1.00 | 6 | 16 | 0 | 5 | 0 |
| pop_center | 0 | 1.00 | 20 | 53 | 0 | 3 | 0 |
| marital_status | 7 | 1.00 | 7 | 21 | 0 | 6 | 0 |
| aboriginal | 3855 | 0.81 | 2 | 10 | 0 | 3 | 0 |
| vis_minority | 140 | 0.99 | 10 | 22 | 0 | 3 | 0 |
| age_immigration | 17225 | 0.16 | 12 | 14 | 0 | 16 | 0 |
| landed_immigrant | 16450 | 0.20 | 2 | 10 | 0 | 3 | 0 |
| citizenship_status | 1143 | 0.94 | 8 | 17 | 0 | 3 | 0 |
| education | 341 | 0.98 | 28 | 60 | 0 | 7 | 0 |
| own_rent | 120 | 0.99 | 10 | 59 | 0 | 3 | 0 |
| living_arrangement | 0 | 1.00 | 5 | 51 | 0 | 12 | 0 |
| hh_type | 76 | 1.00 | 5 | 40 | 0 | 5 | 0 |
| partner_birth_country | 7697 | 0.63 | 6 | 22 | 0 | 3 | 0 |
| partner_birth_province | 7883 | 0.62 | 6 | 39 | 0 | 12 | 0 |
| partner_vis_minority | 7719 | 0.63 | 10 | 22 | 0 | 3 | 0 |
| partner_sex | 20407 | 0.01 | 4 | 10 | 0 | 3 | 0 |
| partner_education | 8259 | 0.60 | 28 | 60 | 0 | 7 | 0 |
| average_hours_worked | 7166 | 0.65 | 6 | 19 | 0 | 6 | 0 |
| worked_last_week | 23 | 1.00 | 2 | 10 | 0 | 3 | 0 |
| partner_main_activity | 7907 | 0.62 | 5 | 51 | 0 | 10 | 0 |
| selfRated_health | 99 | 1.00 | 4 | 10 | 0 | 6 | 0 |
| selfRated_mental_health | 106 | 0.99 | 4 | 10 | 0 | 6 | 0 |
| religion_has_affiliation | 282 | 0.99 | 10 | 25 | 0 | 3 | 0 |
| religion_importance | 253 | 0.99 | 10 | 20 | 0 | 5 | 0 |
| language_home | 448 | 0.98 | 6 | 41 | 0 | 8 | 0 |
| language_knowledge | 105 | 0.99 | 10 | 26 | 0 | 5 | 0 |
| income_family | 0 | 1.00 | 17 | 21 | 0 | 6 | 0 |
| income_respondent | 0 | 1.00 | 17 | 21 | 0 | 6 | 0 |
| occupation | 7297 | 0.65 | 9 | 59 | 0 | 11 | 0 |
| childcare_regular | 18756 | 0.09 | 10 | 35 | 0 | 3 | 0 |
| childcare_type | 19365 | 0.06 | 14 | 38 | 0 | 6 | 0 |
| childcare_monthly_cost | 19962 | 0.03 | 2 | 18 | 0 | 7 | 0 |
| ever_fathered_child | 13604 | 0.34 | 2 | 10 | 0 | 3 | 0 |
| ever_given_birth | 12769 | 0.38 | 2 | 10 | 0 | 3 | 0 |
| number_of_current_union | 18600 | 0.10 | 11 | 21 | 0 | 4 | 0 |
| lives_with_partner | 0 | 1.00 | 2 | 3 | 0 | 2 | 0 |
| children_in_household | 0 | 1.00 | 8 | 22 | 0 | 4 | 0 |
| has_grandchildren | 4 | 1.00 | 2 | 3 | 0 | 2 | 0 |
| grandparents_still_living | 9 | 1.00 | 2 | 10 | 0 | 3 | 0 |
| ever_married | 5 | 1.00 | 2 | 10 | 0 | 3 | 0 |
| current_marriage_is_first | 10416 | 0.49 | 2 | 10 | 0 | 3 | 0 |
| religion_participation | 199 | 0.99 | 10 | 23 | 0 | 6 | 0 |
| partner_location_residence | 18978 | 0.08 | 10 | 36 | 0 | 4 | 0 |

| skim_variable | n_missing | complete_rate | min | max | empty | n_unique | whitespace |
|-----------------------------|-----------|---------------|-----|-----|-------|----------|------------|
| full_part_time_work | 18852 | 0.08 | 9 | 25 | 0 | 3 | 0 |
| time_off_work_birth | 18855 | 0.08 | 2 | 10 | 0 | 3 | 0 |
| reason_no_time_off_birth | 20283 | 0.02 | 5 | 48 | 0 | 10 | 0 |
| returned_same_job | 19451 | 0.06 | 2 | 3 | 0 | 2 | 0 |
| satisfied_time_children | 19691 | 0.04 | 9 | 17 | 0 | 5 | 0 |
| provide_or_receive_fin_supp | 19578 | 0.05 | 10 | 35 | 0 | 5 | 0 |
| fin_supp_agreement | 19937 | 0.03 | 5 | 60 | 0 | 5 | 0 |
| future_children_intention | 13438 | 0.35 | 6 | 18 | 0 | 6 | 0 |
| age_diff | 10430 | 0.49 | 10 | 42 | 0 | 16 | 0 |
| educ_cat | 341 | 0.98 | 8 | 26 | 0 | 6 | 0 |
| partner_educ_cat | 8259 | 0.60 | 8 | 26 | 0 | 6 | 0 |
| has_bachelor_or_higher | 341 | 0.98 | 2 | 3 | 0 | 2 | 0 |

Variable type: logical

| skim_variable | n_missing | complete_rate | mean | count |
|---------------|-----------|---------------|------|-------|
| main_activity | 20602 | 0 | NaN | : |

Variable type: numeric

| skim_variable | n_missing | complete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 |
|----------------------------------|-----------|---------------|----------|---------|------|---------|---------|---------|---------|
| caseid | 0 | 1.00 | 10301.50 | 5947.43 | 1.0 | 5151.25 | 10301.5 | 15451.5 | 15451.5 |
| age | 0 | 1.00 | 52.19 | 17.75 | 15.0 | 37.30 | 54.2 | 60.0 | 60.0 |
| age_first_child | 6835 | 0.67 | 30.57 | 17.10 | 0.0 | 15.00 | 32.0 | 40.0 | 40.0 |
| age_youngest_child_under_6 | 18488 | 0.10 | 2.41 | 1.60 | 0.0 | 1.00 | 2.0 | 4.0 | 4.0 |
| total_children | 19 | 1.00 | 1.68 | 1.49 | 0.0 | 0.00 | 2.0 | 3.0 | 3.0 |
| age_start_relationship | 18566 | 0.10 | 33.63 | 11.20 | 18.0 | 25.00 | 30.5 | 40.0 | 40.0 |
| age_at_first_marriage | 15248 | 0.26 | 24.10 | 5.41 | 15.0 | 20.50 | 22.8 | 26.0 | 26.0 |
| age_at_first_birth | 7865 | 0.62 | 26.86 | 5.42 | 18.0 | 22.80 | 26.4 | 30.0 | 30.0 |
| distance_between_houses | 19476 | 0.05 | 17.13 | 18.18 | 0.0 | 4.00 | 10.0 | 20.0 | 20.0 |
| age_youngest_child_returned_work | 19466 | 0.06 | 6.59 | 6.17 | 0.2 | 0.50 | 6.0 | 10.0 | 10.0 |
| feelings_life | 271 | 0.99 | 8.09 | 1.65 | 0.0 | 7.00 | 8.0 | 9.0 | 9.0 |
| hh_size | 0 | 1.00 | 2.35 | 1.26 | 1.0 | 1.00 | 2.0 | 3.0 | 3.0 |
| number_total_children_intention | 12202 | 0.41 | 0.90 | 1.18 | 0.0 | 0.00 | 0.0 | 1.0 | 1.0 |
| number_marriages | 0 | 1.00 | 0.80 | 0.62 | 0.0 | 0.00 | 1.0 | 1.0 | 1.0 |
| fin_supp_child_supp | 20057 | 0.03 | 0.77 | 0.42 | 0.0 | 1.00 | 1.0 | 1.0 | 1.0 |
| fin_supp_child_exp | 20057 | 0.03 | 0.34 | 0.47 | 0.0 | 0.00 | 0.0 | 1.0 | 1.0 |
| fin_supp_lump | 20057 | 0.03 | 0.06 | 0.23 | 0.0 | 0.00 | 0.0 | 0.0 | 0.0 |
| fin_supp_other | 20057 | 0.03 | 0.06 | 0.23 | 0.0 | 0.00 | 0.0 | 0.0 | 0.0 |
| is_male | 0 | 1.00 | 0.46 | 0.50 | 0.0 | 0.00 | 0.0 | 1.0 | 1.0 |
| number_total_children_known | 0 | 1.00 | 0.41 | 0.49 | 0.0 | 0.00 | 0.0 | 1.0 | 1.0 |
| age_group | 0 | 1.00 | 49.88 | 17.97 | 15.0 | 35.00 | 50.0 | 60.0 | 60.0 |

Calculating the number of rows, columns

Use `nrow()`, `ncol()` and `dim()`

```
nrow(gss)
```

```
## [1] 20602
```

```
ncol(gss)
```

```
## [1] 85
```

```
dim(gss)
```

```
## [1] 20602    85
```

```
# or can pipe
```

```
gss %>%  
  filter(place_birth_canada == "Born in Canada") %>%  
  nrow()
```

```
## [1] 16355
```

Handling Categorical Data

The group_by function

The `group_by` function allows you to get key summary statistics by group (levels of a categorical variable). Use in combination with `summarize` etc that we learnt last week.

e.g. mean life expectancy by region in 2017

```
country_ind %>%  
  filter(year == 2017) %>%  
  group_by(region) %>%  
  summarize(mean_le = mean(life_expectancy))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 10 x 2  
##   region                mean_le  
##   <chr>                <dbl>  
## 1 Caucasus and Central Asia    75.2  
## 2 Developed regions           82.3  
## 3 Eastern Asia                79.4  
## 4 Latin America and Caribbean 77.6  
## 5 Northern Africa             76.9  
## 6 Oceania                    71.5  
## 7 South-eastern Asia          76.1  
## 8 Southern Asia               73.2  
## 9 Sub-Saharan Africa          64.3  
## 10 Western Asia               77.2
```

e.g. mean age and standard deviation by marital status in GSS

```
gss %>%
  group_by(marital_status) %>%
  summarize(mean_age = mean(age),
            sd_age = sd(age)) %>%
  arrange(mean_age)

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 7 x 3
##   marital_status      mean_age sd_age
##   <chr>              <dbl> <dbl>
## 1 Single, never married    38.1  17.2
## 2 Living common-law      44.6  14.5
## 3 Separated              54.5  13.7
## 4 Married                54.9  14.8
## 5 Divorced               61.0  11.4
## 6 <NA>                  65.8  12.9
## 7 Widowed               73.0   8.47
```

Note that the above table shows the mean and sd of age for when marital status is missing (NA). We may want to remove those. To do this, use the `is.na` function in combination with the `!` (which means “not”)

```
gss %>%
  filter(!is.na(marital_status)) %>%
  group_by(marital_status) %>%
  summarize(mean_age = mean(age),
            sd_age = sd(age)) %>%
  arrange(mean_age)

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 6 x 3
##   marital_status      mean_age sd_age
##   <chr>              <dbl> <dbl>
## 1 Single, never married    38.1  17.2
## 2 Living common-law      44.6  14.5
## 3 Separated              54.5  13.7
## 4 Married                54.9  14.8
## 5 Divorced               61.0  11.4
## 6 Widowed               73.0   8.47
```

Note dealing with missing data is a significant part of data analysis. While in some analysis we decide to exclude missing observations, take a moment and think about why some observations may be missing.

Calculating the correlation coefficient

To calculate the correlation coefficient between two quantitative (numerical/continuous) variables, e.g. age and age at first marriage, use the `summarize` function. Notice that we need to remove rows with any NA values before doing the calculation. We can do this using `drop_na()`

```
gss %>%
  select(age, age_at_first_marriage) %>%
  drop_na() %>%
  summarise(correlation = cor(age, age_at_first_marriage))
```

```
## # A tibble: 1 x 1
##   correlation
##   <dbl>
## 1      -0.154
```

Counts and proportions

Counting the number of observations

Often we would like to include counts of observations in particular groups. To do this, use the `tally()` or `count()` function.

e.g. the number of people by province of residence in the GSS

```
gss %>%
  group_by(province) %>%
  tally()
```

```
## # A tibble: 10 x 2
##   province      n
##   <chr>    <int>
## 1 Alberta      1728
## 2 British Columbia 2522
## 3 Manitoba      1192
## 4 New Brunswick  1337
## 5 Newfoundland and Labrador 1094
## 6 Nova Scotia    1425
## 7 Ontario      5621
## 8 Prince Edward Island 708
## 9 Quebec      3822
## 10 Saskatchewan  1153
```

equivalent:

```
gss %>%
  count(province)
```

```
## # A tibble: 10 x 2
##   province      n
##   <chr>    <int>
## 1 Alberta      1728
## 2 British Columbia 2522
## 3 Manitoba      1192
## 4 New Brunswick  1337
## 5 Newfoundland and Labrador 1094
```



```
## 6 Nova Scotia          1425
## 7 Ontario              5621
## 8 Prince Edward Island   708
## 9 Quebec               3822
## 10 Saskatchewan         1153
```

Getting the proportion in each group

Also often useful to get proportion of total in each group:

```
gss %>%
  group_by(province) %>%
  tally() %>%
  mutate(prop = n / sum(n))
```

```
## # A tibble: 10 x 3
##   province          n   prop
##   <chr>          <int> <dbl>
## 1 Alberta        1728 0.0839
## 2 British Columbia 2522 0.122
## 3 Manitoba        1192 0.0579
## 4 New Brunswick    1337 0.0649
## 5 Newfoundland and Labrador 1094 0.0531
## 6 Nova Scotia      1425 0.0692
## 7 Ontario         5621 0.273
## 8 Prince Edward Island   708 0.0344
## 9 Quebec         3822 0.186
## 10 Saskatchewan      1153 0.0560
```

equivalent

```
gss %>%
  count(province) %>%
  mutate(prop = n / sum(n))
```

```
## # A tibble: 10 x 3
##   province          n   prop
##   <chr>          <int> <dbl>
## 1 Alberta        1728 0.0839
## 2 British Columbia 2522 0.122
## 3 Manitoba        1192 0.0579
## 4 New Brunswick    1337 0.0649
## 5 Newfoundland and Labrador 1094 0.0531
## 6 Nova Scotia      1425 0.0692
## 7 Ontario         5621 0.273
## 8 Prince Edward Island   708 0.0344
## 9 Quebec         3822 0.186
## 10 Saskatchewan      1153 0.0560
```

In-class exercise

1. What proportion of age of first marriage is missing?

2. What are the proportion of individuals have worked last week (`worked_last_week`)? What proportion of this variable is missing?
3. Within non-missing individuals who have worked last week, how many and what proportion worked full-time (`full_part_time_work`)?

ggplot

`ggplot` is a powerful visualization package. It provides many options to make beautiful graphs, maps, plots of all sort. Each example we look at today.

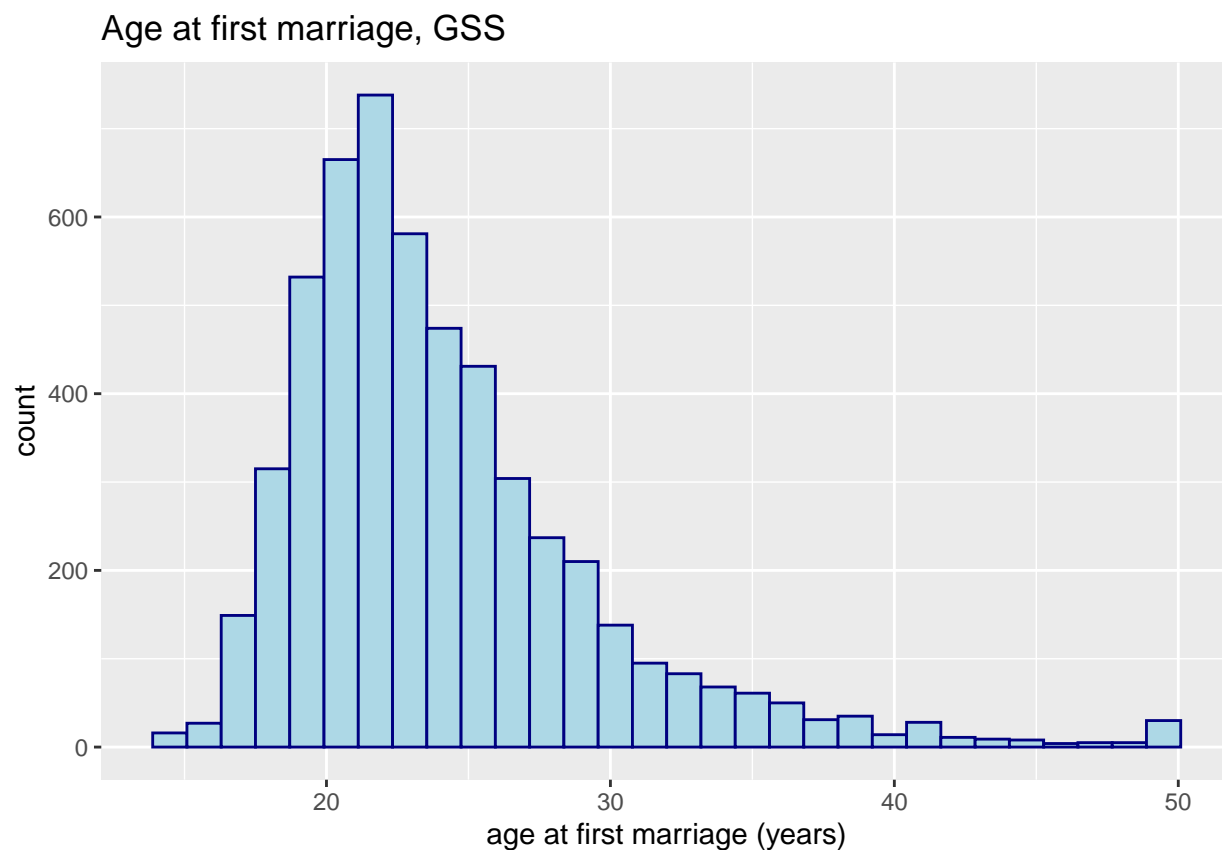
Histograms

Note for histograms, bar chats, box plots, `fill` is the main color choice (`color` changes the outline)

```
ggplot(data = gss, aes(age_at_first_marriage)) +  
  geom_histogram(fill = "lightblue", color = "navy") +  
  ggtitle("Age at first marriage, GSS") +  
  xlab("age at first marriage (years)")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 15248 rows containing non-finite values (stat_bin).
```



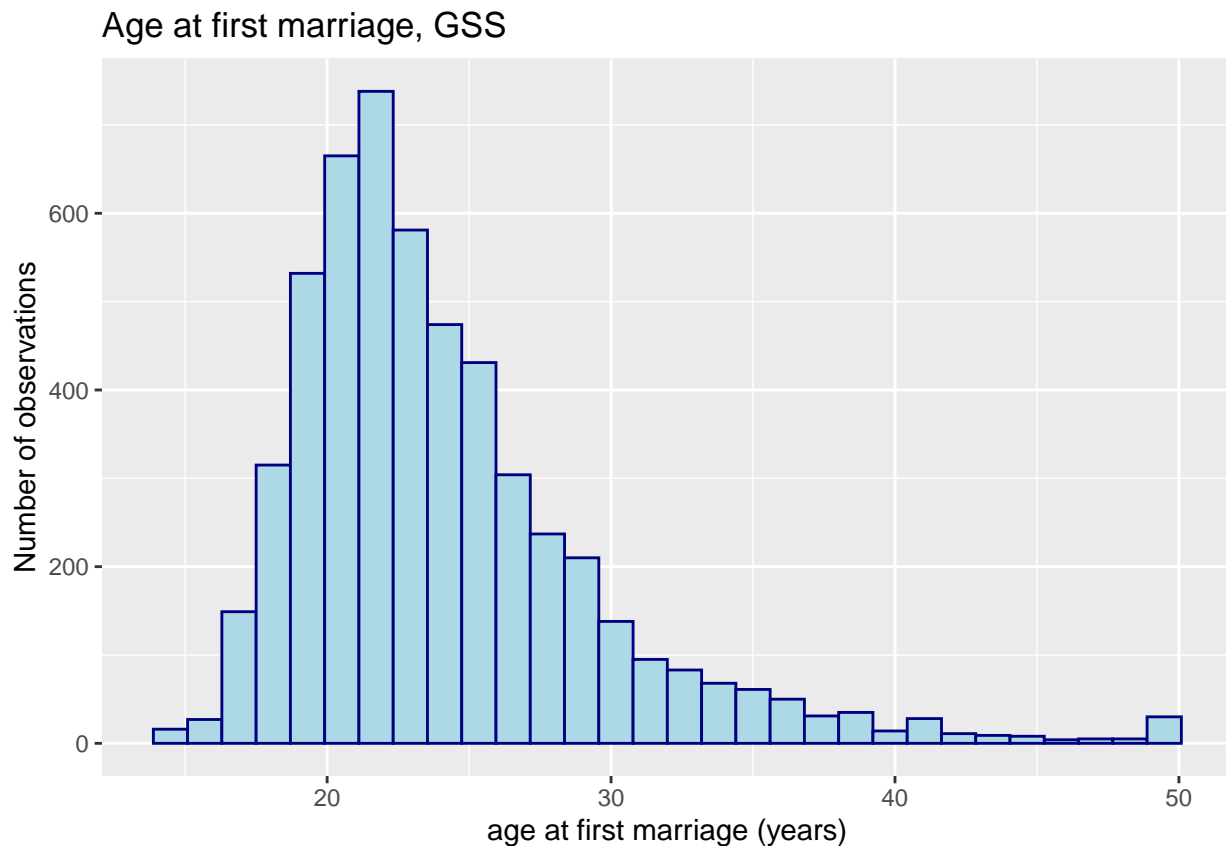
Note that you can also save the plot as an object and then print it

```
my_plot <- ggplot(data = gss, aes(age_at_first_marriage)) +
  geom_histogram(fill = "lightblue", color = "navy")+
  ggtitle("Age at first marriage, GSS") +
  xlab("age at first marriage (years)")

# print
my_plot + ylab("Number of observations")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 15248 rows containing non-finite values (stat_bin).
```

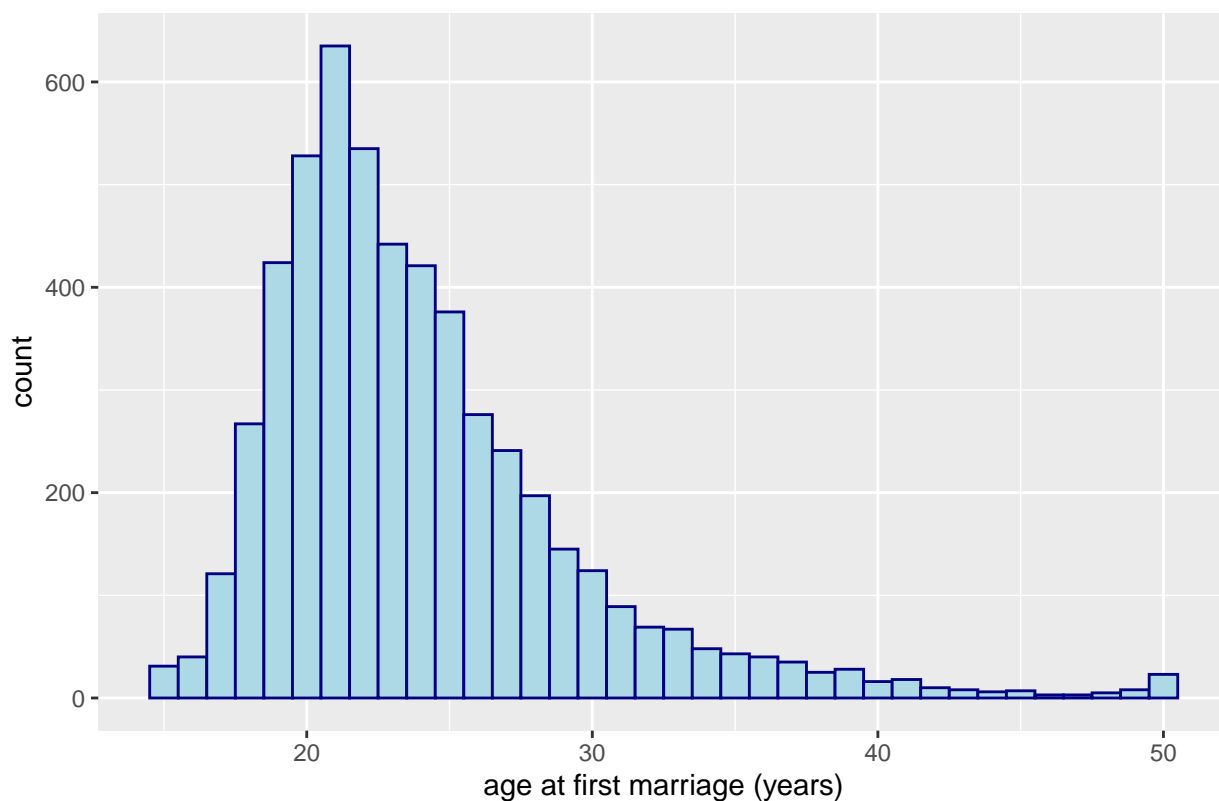


Histograms select a `binwidth` or section of the data and then count how many of the observations fall within that. Histograms look different depending on the size of the bins. You can also supply the number of bins that you want to create.

```
ggplot(data = gss, aes(age_at_first_marriage)) +
  geom_histogram(fill = "lightblue", color = "navy", binwidth = 1) +
  ggtitle("Age at first marriage, GSS") +
  xlab("age at first marriage (years)")
```

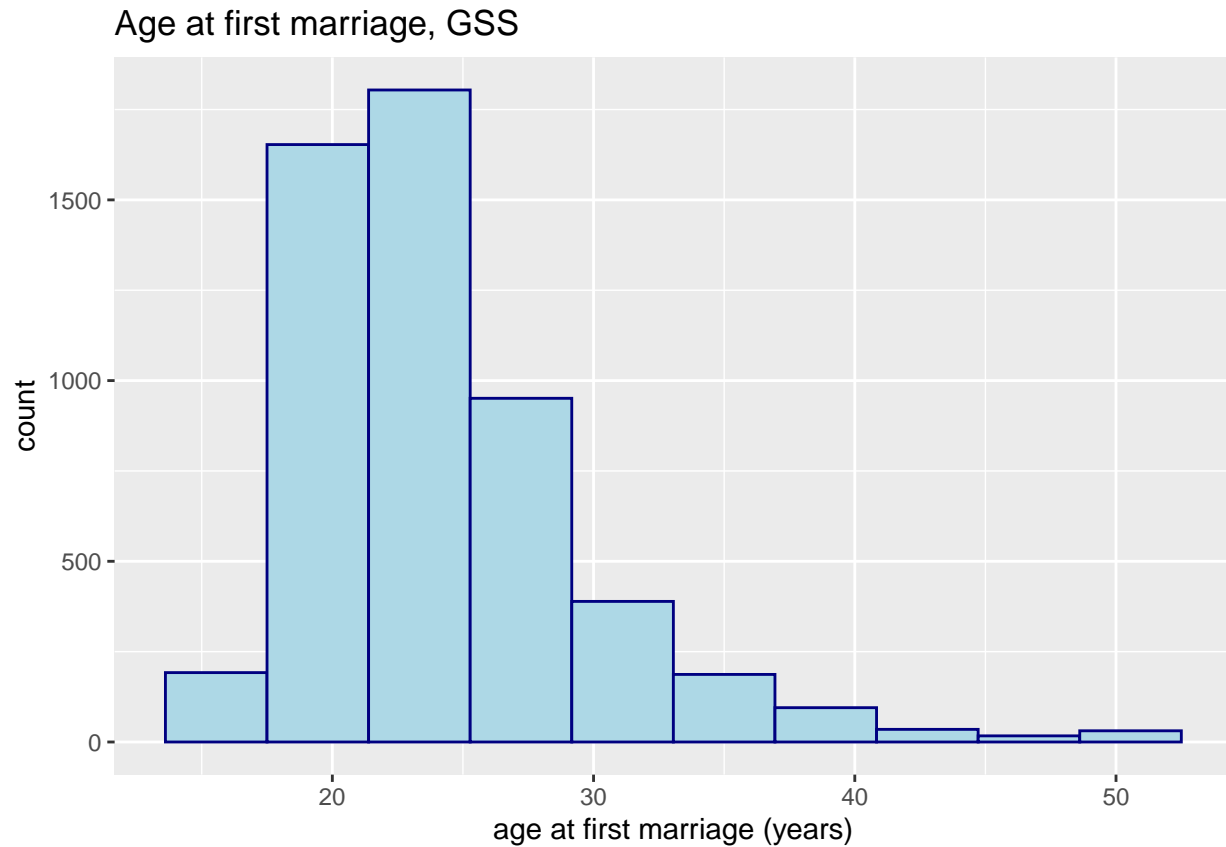
```
## Warning: Removed 15248 rows containing non-finite values (stat_bin).
```

Age at first marriage, GSS



```
ggplot(data = gss, aes(age_at_first_marriage)) +  
  geom_histogram(fill = "lightblue", color = "navy", bins = 10)+  
  ggtitle("Age at first marriage, GSS") +  
  xlab("age at first marriage (years)")
```

Warning: Removed 15248 rows containing non-finite values (stat_bin).

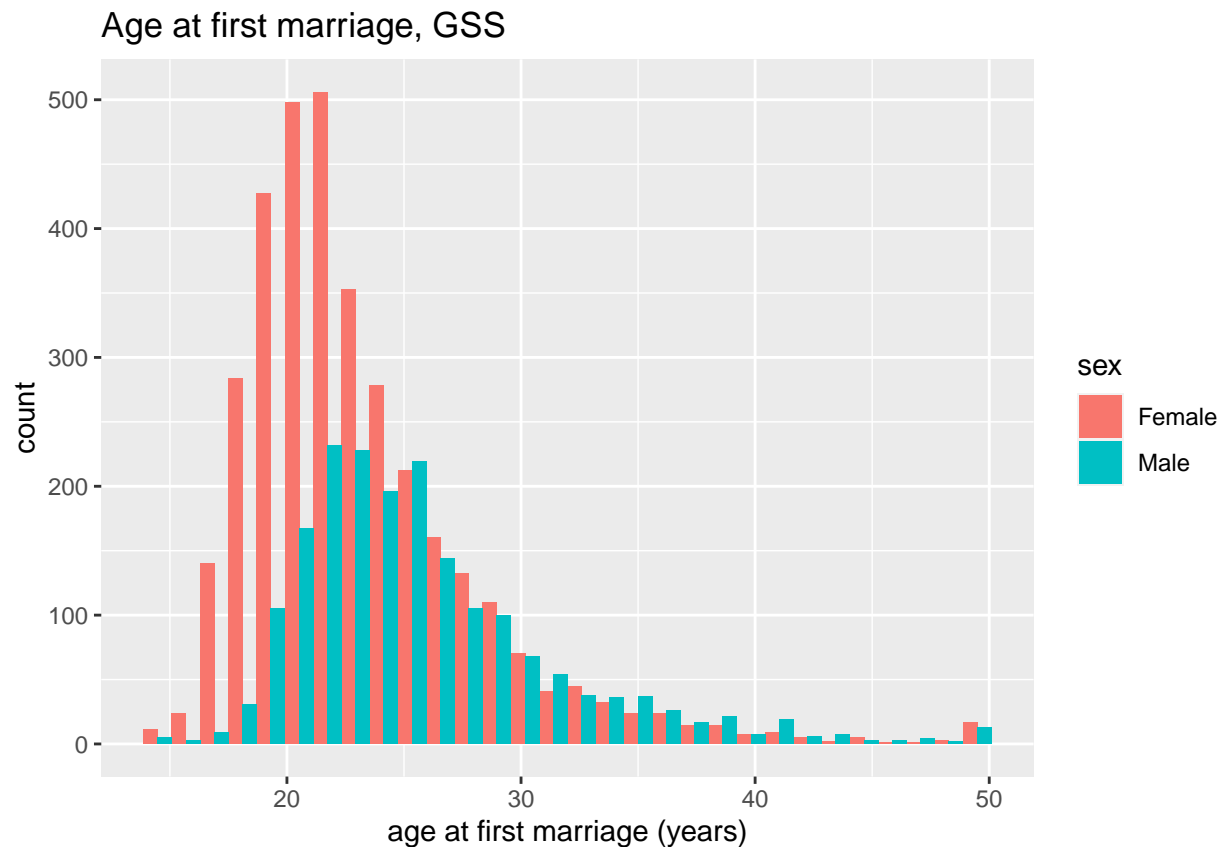


We can also plot by another variable to compare the plots by the categories of the variable. For example, we look at plots by sex:

```
ggplot(data = gss, aes(age_at_first_marriage, fill = sex)) +  
  geom_histogram(position = 'dodge') +  
  ggtitle("Age at first marriage, GSS") +  
  xlab("age at first marriage (years)")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 15248 rows containing non-finite values (stat_bin).
```



Bar charts

Let's plot the proportion of respondents by province as a bar chart. First save the proportions as a new data frame

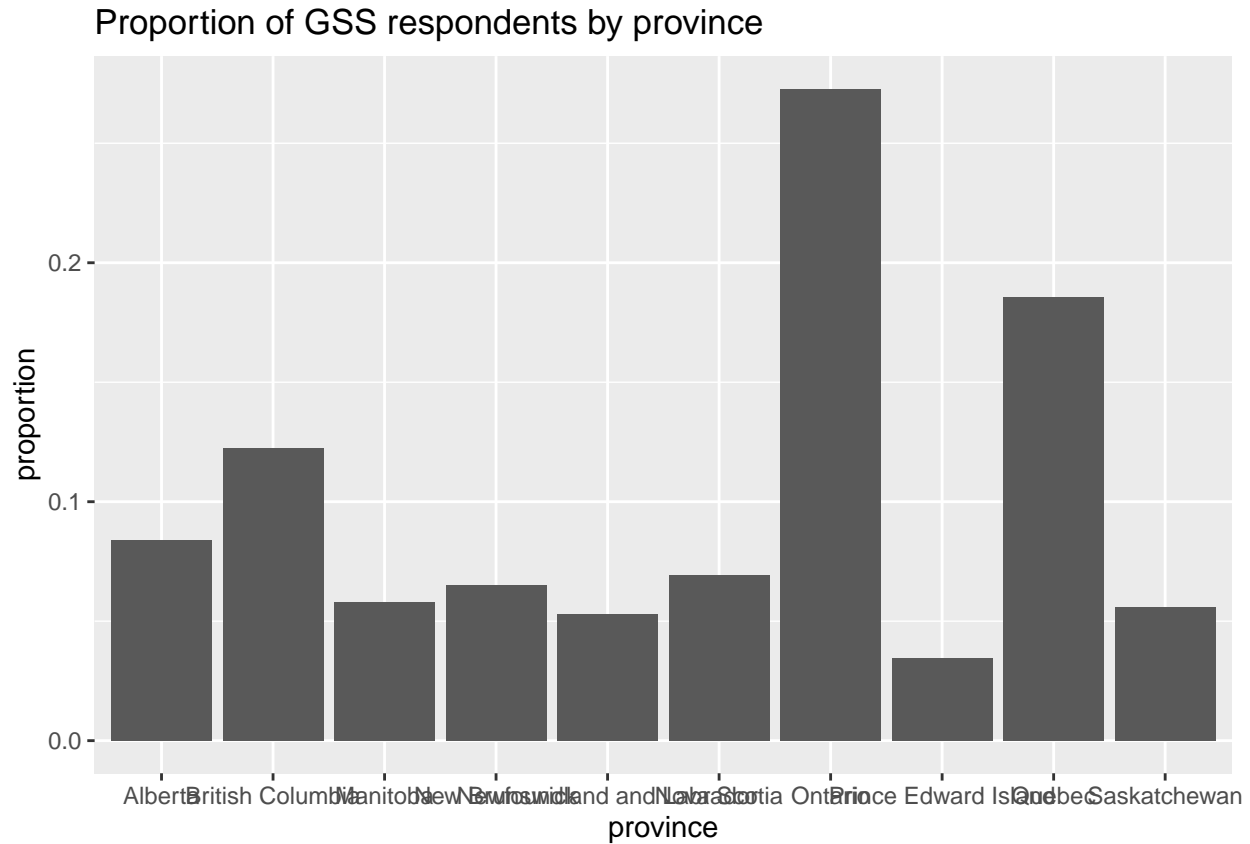
```
resp_by_prov <- gss %>%
  group_by(province) %>%
  tally() %>%
  mutate(prop = n / sum(n))
```

```
resp_by_prov
```

```
## # A tibble: 10 x 3
##   province      n  prop
##   <chr>    <int> <dbl>
## 1 Alberta    1728 0.0839
## 2 British Columbia 2522 0.122
## 3 Manitoba   1192 0.0579
## 4 New Brunswick 1337 0.0649
## 5 Newfoundland and Labrador 1094 0.0531
## 6 Nova Scotia 1425 0.0692
## 7 Ontario   5621 0.273
## 8 Prince Edward Island 708 0.0344
## 9 Quebec    3822 0.186
## 10 Saskatchewan 1153 0.0560
```

Now plot

```
ggplot(data = resp_by_prov, aes(x = province, y = prop)) +  
  geom_bar(stat = "identity") +  
  ylab("proportion") +  
  ggtitle("Proportion of GSS respondents by province")
```

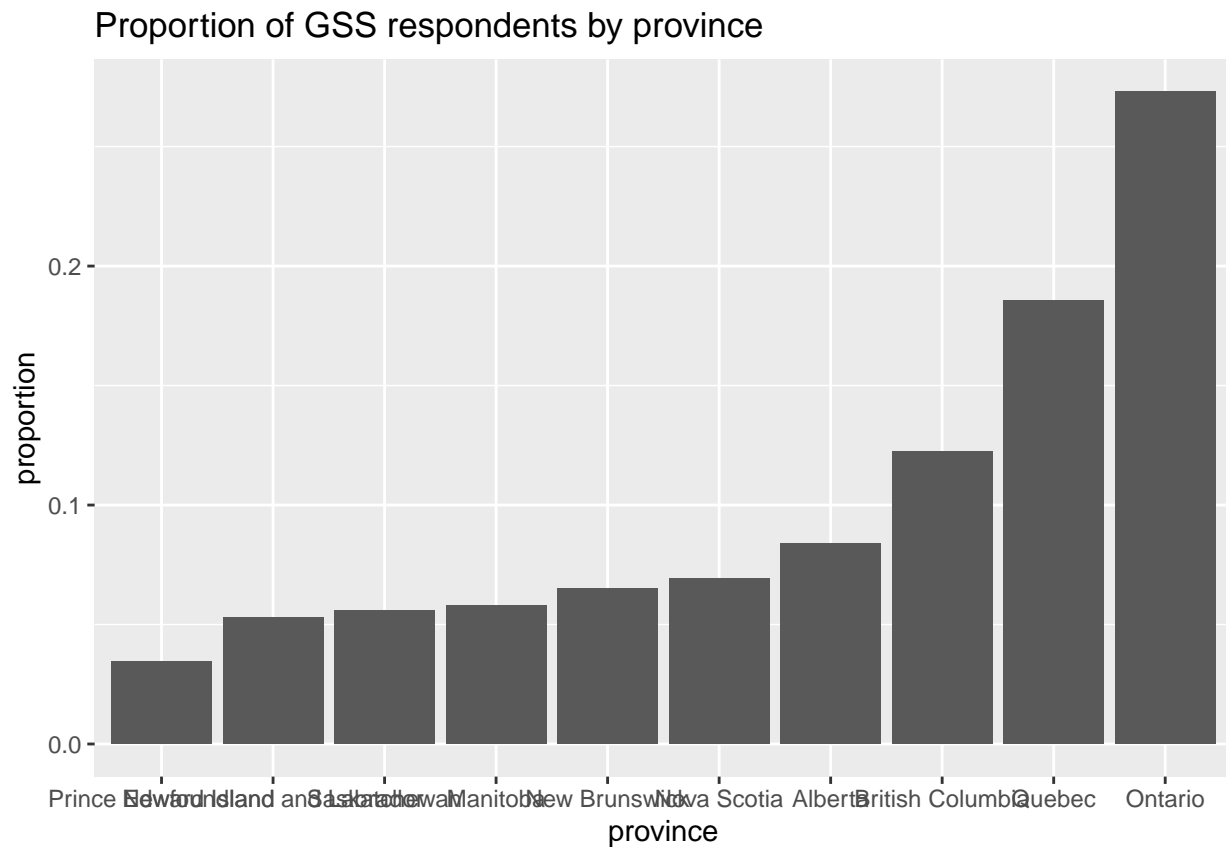


There are a few things here that would be nice to fix. Firstly, the categories are ordered alphabetically, which is the default. It would be better visually to order by proportion. We can do this using the `fct_reorder` function to alter (mutate) the province variable.

```
resp_by_prov <- resp_by_prov %>%  
  mutate(province = fct_reorder(province, prop)) # order by proportion
```

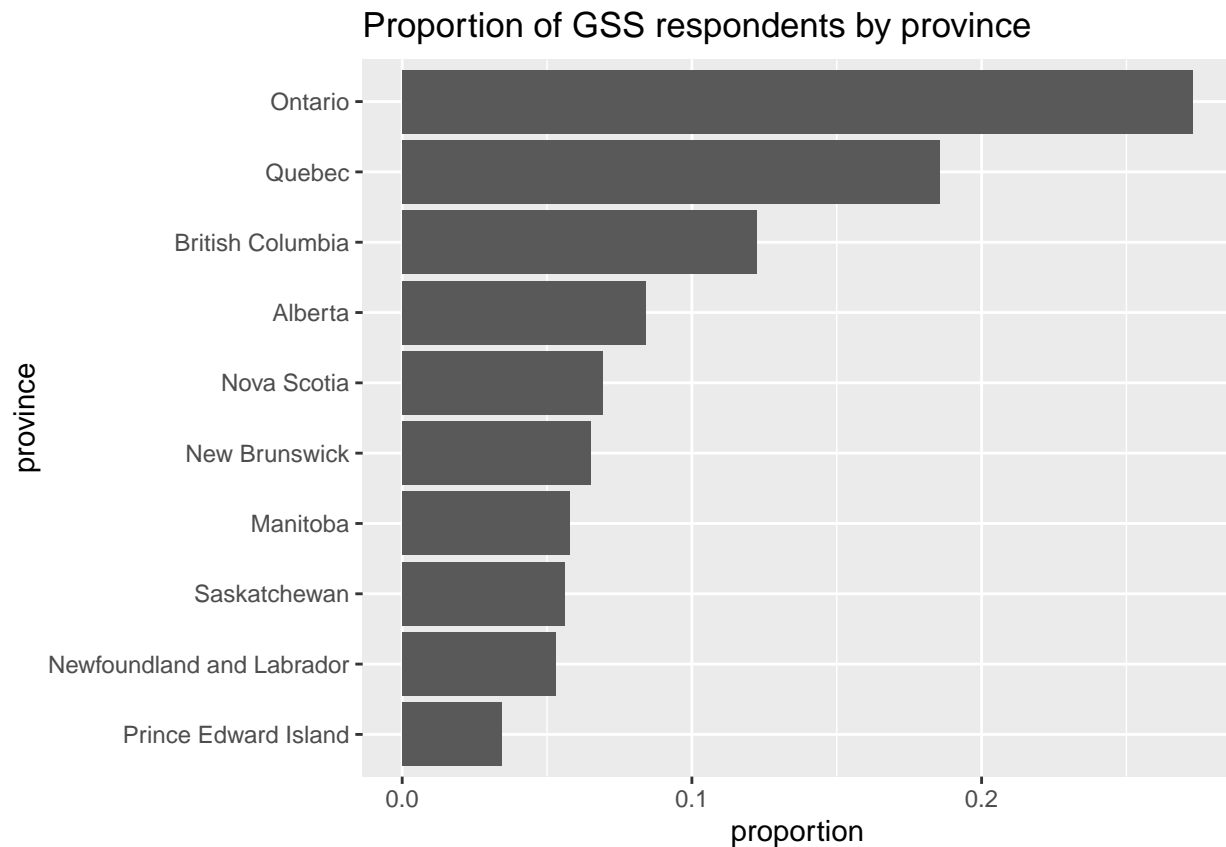
Now try plotting again.

```
ggplot(data = resp_by_prov, aes(x = province, y = prop)) +  
  geom_bar(stat = "identity") +  
  ylab("proportion") +  
  ggtitle("Proportion of GSS respondents by province")
```



To improve readability, could change to horizontal bar chart.

```
ggplot(data = resp_by_prov, aes(x = province, y = prop)) +  
  geom_bar(stat = "identity") +  
  ylab("proportion") +  
  ggtitle("Proportion of GSS respondents by province") +  
  coord_flip()
```

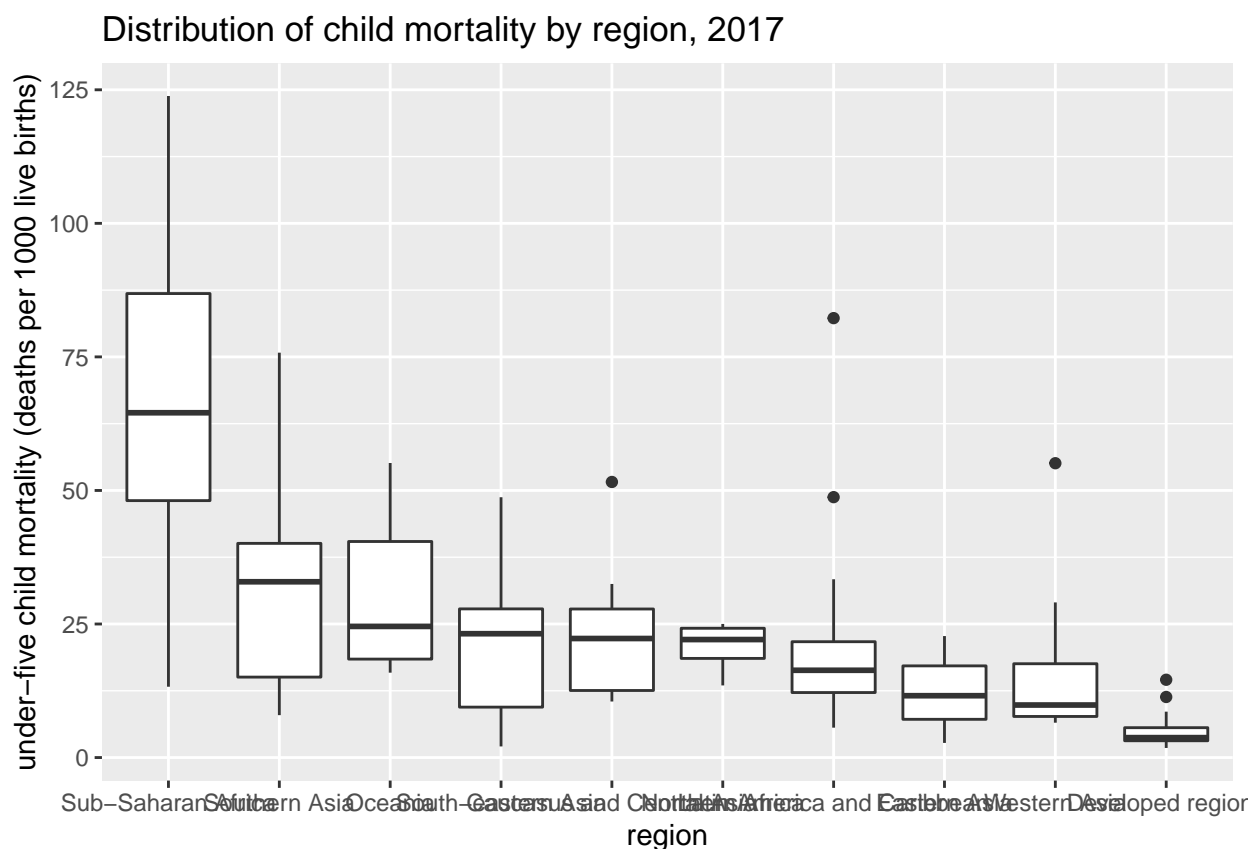



Box plots

Let's use the country indicators dataset here and do boxplots of child mortality in 2017 over regions. Like the bar chart example, best to reorder the regions by the variable we are interested in

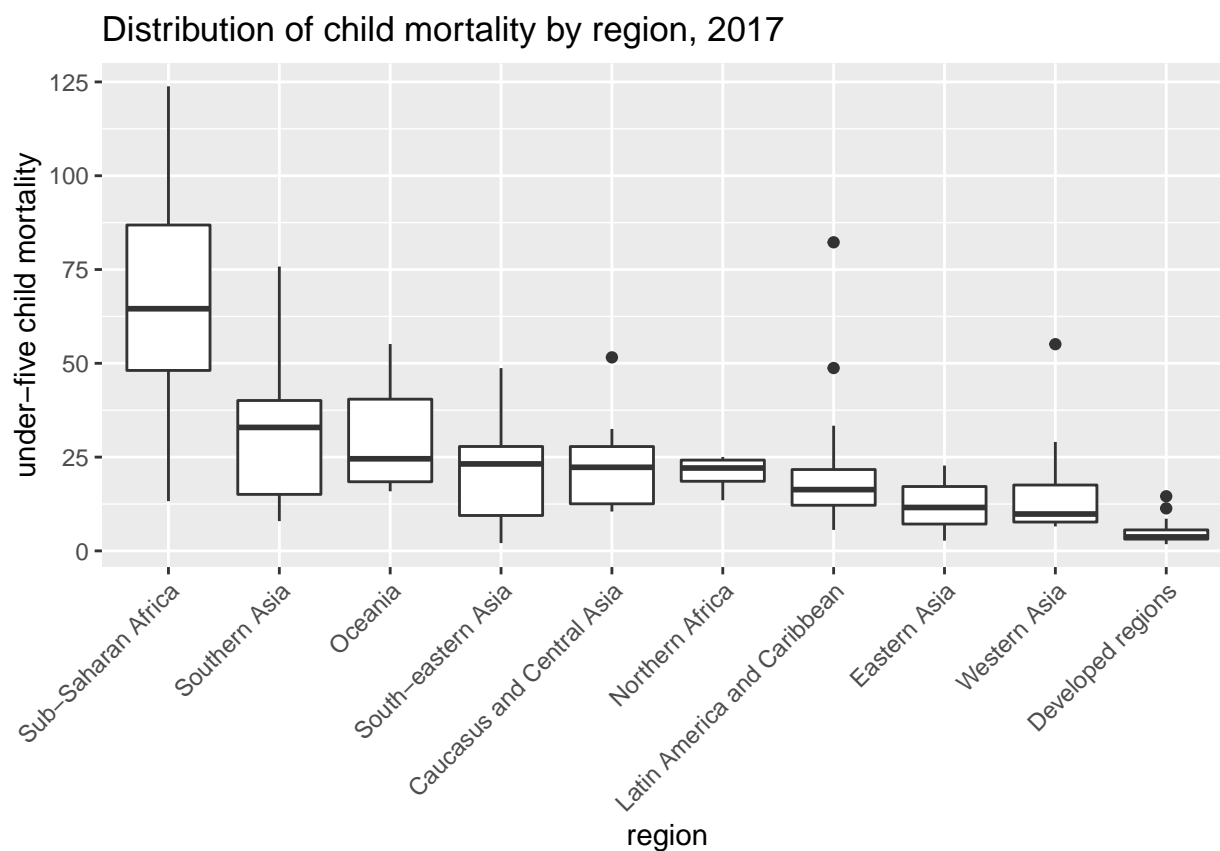
```
country_ind_2017 <- country_ind %>%
  filter(year==2017) %>%
  mutate(region = fct_reorder(region, -child_mort)) # descending order

ggplot(data = country_ind_2017, aes(x = region, y = child_mort)) +
  geom_boxplot() +
  ylab("under-five child mortality (deaths per 1000 live births)") +
  ggtitle("Distribution of child mortality by region, 2017")
```



The labels on the x axis are hard to read. We could do the same as last time (switch to horizontal), or we can change the alignment of the labels:

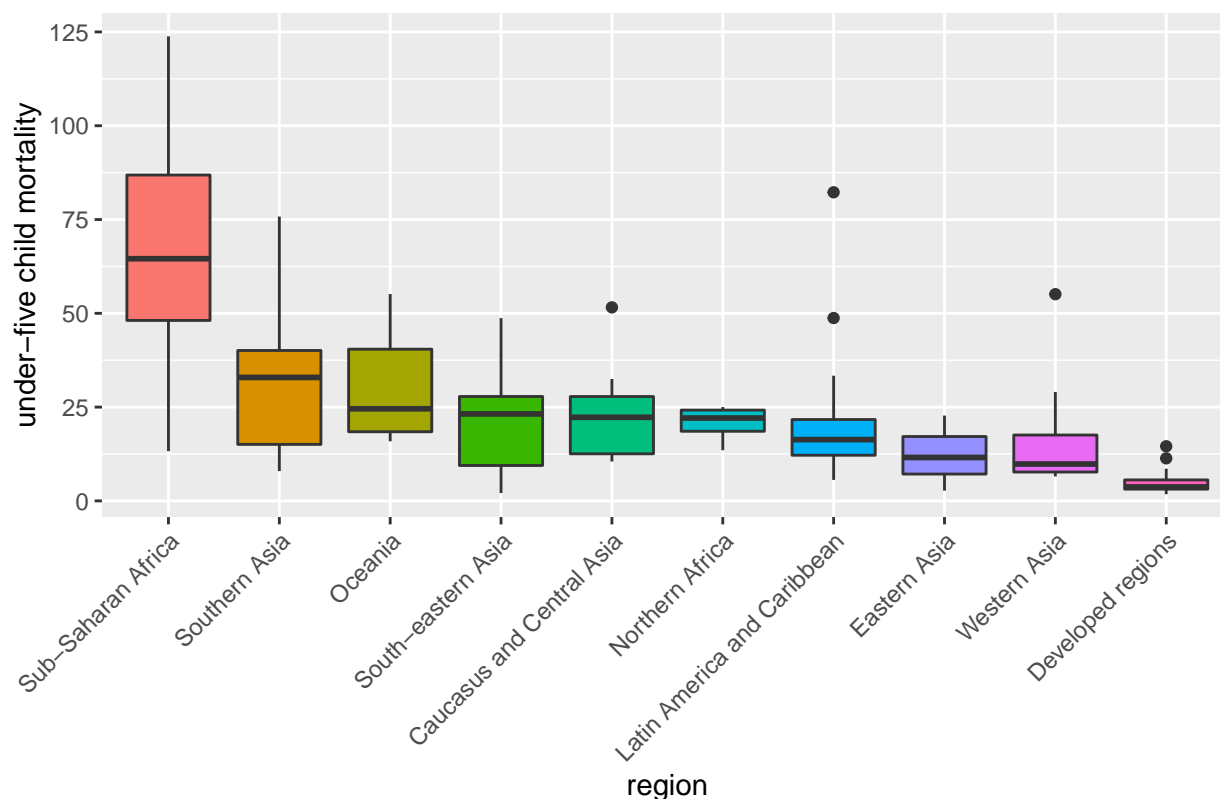
```
ggplot(data = country_ind_2017, aes(x = region, y = child_mort)) +
  geom_boxplot() +
  ylab("under-five child mortality") +
  ggtitle("Distribution of child mortality by region, 2017") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



Note if you want to color the boxes, use `fill`, and then remove the legend (not needed)

```
ggplot(data = country_ind_2017, aes(x = region, y = child_mort, fill = region)) +
  geom_boxplot() +
  ylab("under-five child mortality") +
  ggtitle("Distribution of child mortality by region, 2017") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1) ,
        legend.position = 'none')
```

Distribution of child mortality by region, 2017



Line graphs

Let's look at the mean age at marriage by age of respondent. Firstly, let's make a new variable in the `gss` dataset that groups people into 5-year age groups. Here's the code to do this:

```
age_groups <- seq(15, 80, by = 5)
gss$age_group <- as.numeric(as.character(cut(gss$age,
      breaks= c(age_groups, Inf),
      labels = age_groups,
      right = FALSE)))
```

```
gss %>% select(age, age_group)
```

```
## # A tibble: 20,602 x 2
##   age age_group
##   <dbl>   <dbl>
## 1  52.7     50
## 2  51.1     50
## 3  63.6     60
## 4   80     80
## 5   28     25
## 6   63     60
## 7  58.8     55
## 8   80     80
## 9  63.8     60
```

```
## 10 25.2      25
## # ... with 20,592 more rows
```

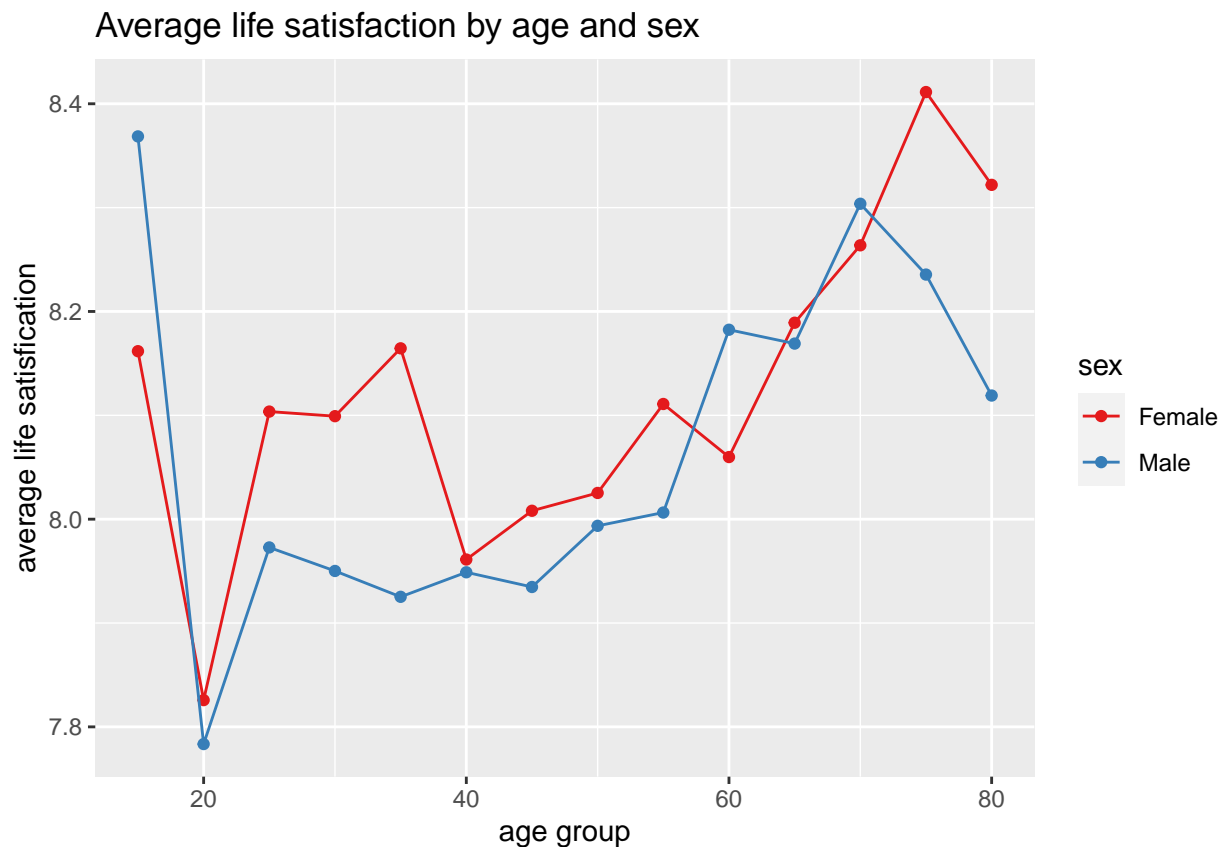
Now let's calculate the average of the 'life satisfaction' variable by age group and sex. This involves a `group_by` by two variables:

```
life_satis_age_sex <- gss %>%
  group_by(age_group, sex) %>%
  summarise(mean_life_satis = mean(feelings_life, na.rm = TRUE))
```

```
## `summarise()` regrouping output by 'age_group' (override with `.groups` argument)
```

Plot as a line chart over age, coloring by sex, for this example we use a different colour palette called "Set1":

```
ggplot(data = life_satis_age_sex, aes(x = age_group, y = mean_life_satis, colour = sex)) +
  geom_point() +
  geom_line() +
  scale_color_brewer(palette = "Set1") + # change the color scheme
  ylab("average life satisfaction") +
  xlab("age group") +
  ggtitle("Average life satisfaction by age and sex")
```

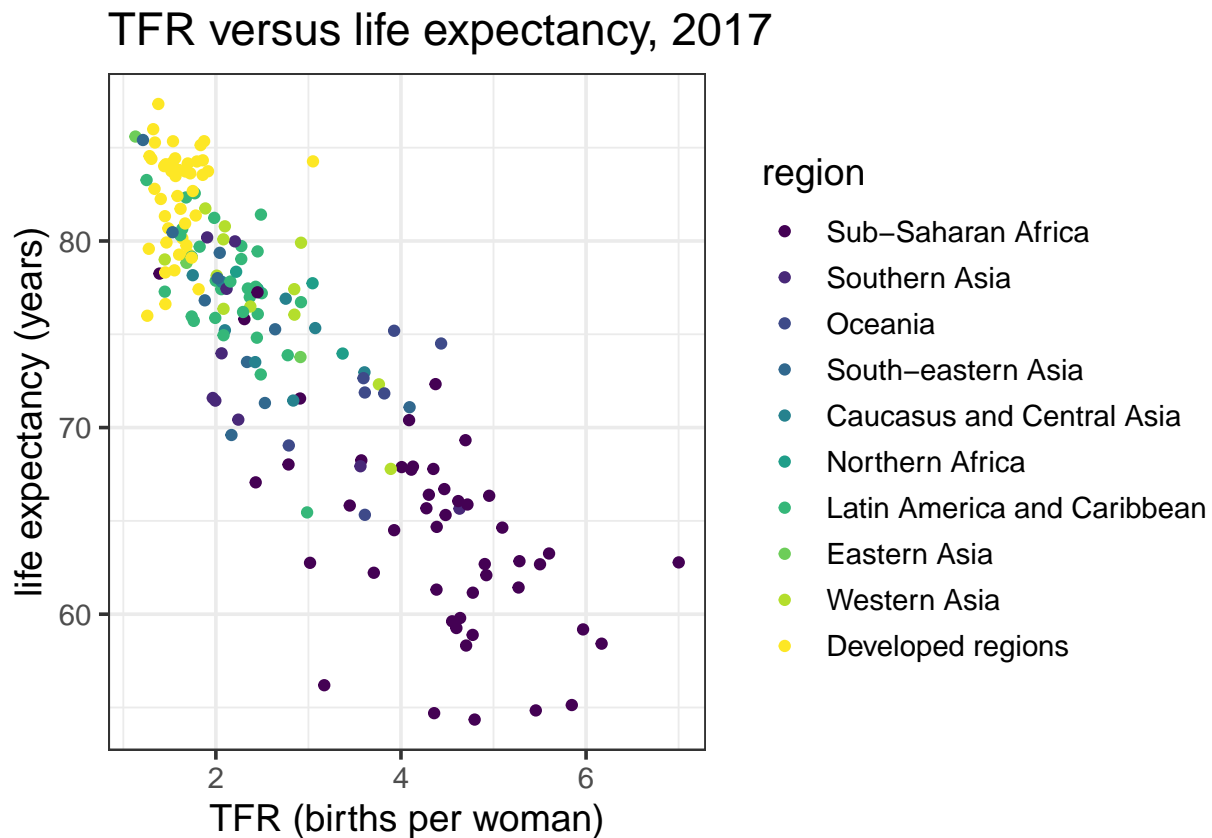


Scatter plots

Let's use the country indicators dataset here. The example in the lecture slides is life expectancy versus TFR. We also used a new colour palette called `viridis`, these colours palettes are designed to be viewable in

black and white as well.

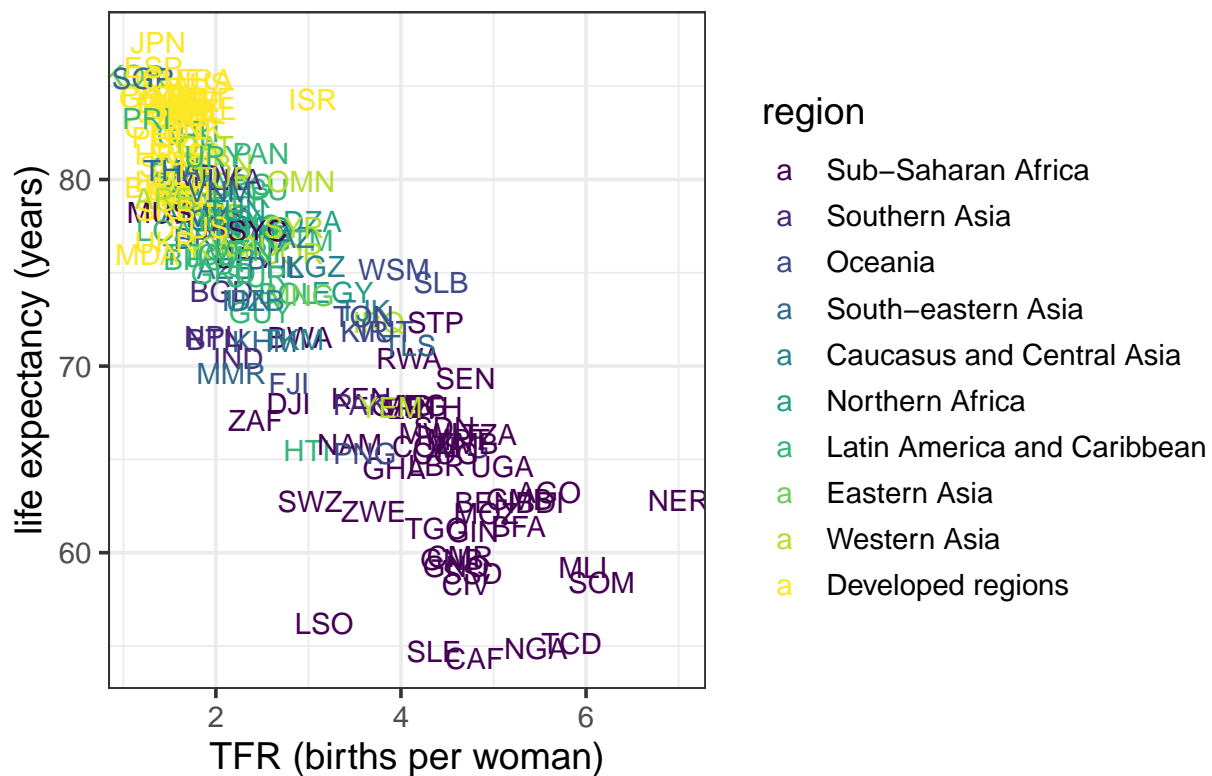
```
ggplot(country_ind_2017, aes(tfr, life_expectancy, color = region,)) +  
  geom_point() +  
  ggtitle("TFR versus life expectancy, 2017")+  
  theme_bw(base_size = 14) +  
  ylab("life expectancy (years)") +  
  xlab("TFR (births per woman)") +  
  scale_color_viridis_d()
```



Instead of dots could have country codes (although becomes hard to read, but easy to see outliers)

```
ggplot(country_ind_2017, aes(tfr, life_expectancy, color = region, label = country_code)) + # adding  
  geom_text() +  
  ggtitle("TFR versus life expectancy, 2017")+  
  theme_bw(base_size = 14)+  
  ylab("life expectancy (years)") +  
  xlab("TFR (births per woman)") +  
  scale_color_viridis_d()
```

TFR versus life expectancy, 2017



Faceting

Changing the color and fills is useful to show one other variable on a graph. For more complicated set-ups, faceting graphs by an additional variable becomes useful.

For example let's go back to plotting a histogram of age at first marriage by sex, but also add in whether or not the respondent was born in Canada. First, look at the unique values of the `place_birth_canada` variable:

```
gss %>%
  select(place_birth_canada) %>%
  unique()
```

```
## # A tibble: 4 x 1
##   place_birth_canada
##   <chr>
## 1 Born in Canada
## 2 Born outside Canada
## 3 <NA>
## 4 Don't know
```

For now, filter the data to only include the first two categories. To do this, use the `%in%` function within filter:

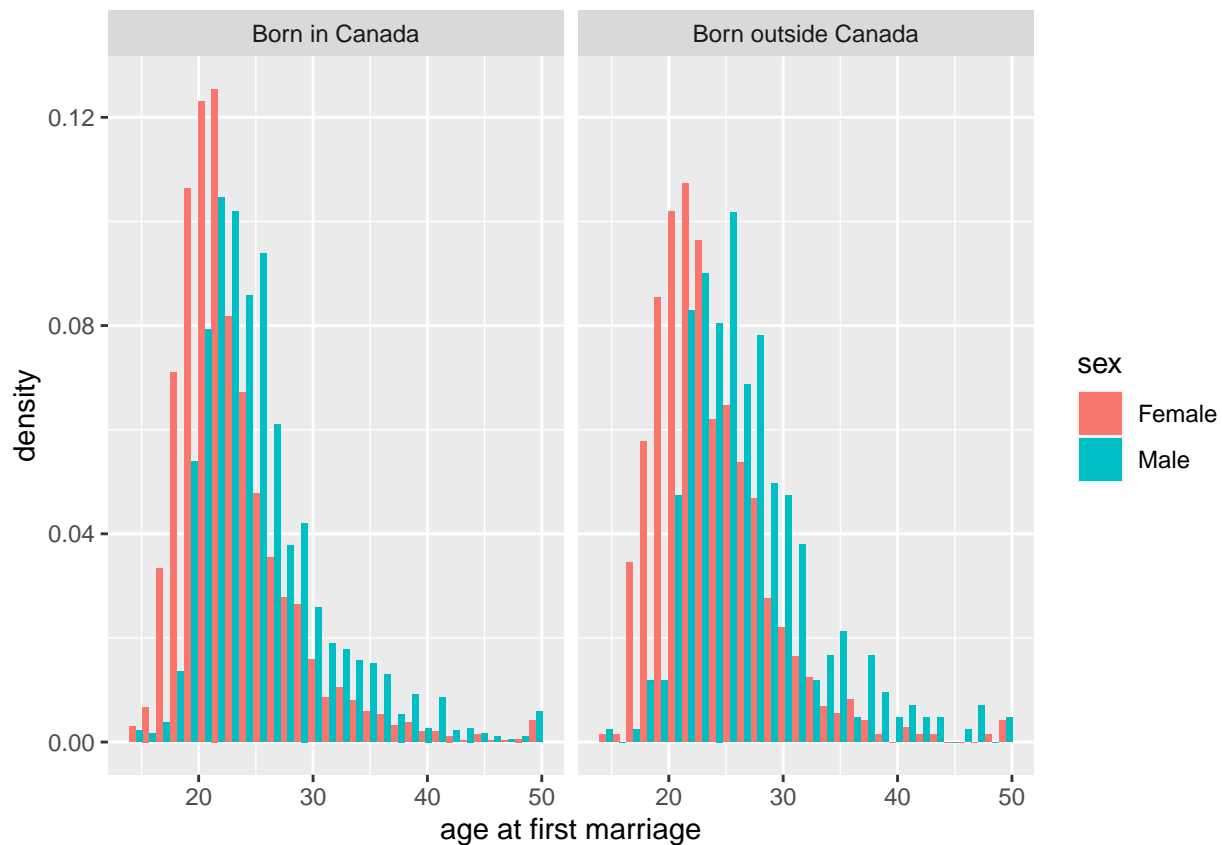
```
gss_subset <- gss %>%
  filter(place_birth_canada %in% c("Born in Canada", "Born outside Canada"))
```

Now plot the histograms as before, but now also facet by place of birth. Note we are plotting the density here.

```
ggplot(data = gss_subset, aes(age_at_first_marriage, fill = sex)) +
  geom_histogram(position = 'dodge', aes(y = ..density..)) +
  facet_wrap(~place_birth_canada) +
  xlab("age at first marriage")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 15137 rows containing non-finite values (stat_bin).
```



Review Questions

1. Using the `country_indicator` dataset, create a scatter plot of GDP over life expectancy by region for the year 2014. Edit the labels, set a title, and make sure the graph is color-coded.
2. Using the GSS dataset, create a bar graph of non-missing values for the province of birth (`'place_birth_province'`) and then arrange the proportions from high to low. Make sure to color code and make all labels are readable.