

SOC6707: Intermediate Data Analysis

Monica Alexander

Week 2: Exploratory Data Analysis and Data Visualization

Overview

What we will cover today:

- ▶ What is EDA and why do we do it?
- ▶ Steps of EDA
- ▶ Data visualization
- ▶ Key types of graphs
 - ▶ another summary measure: correlation coefficient
- ▶ Doing EDA and data viz in R (intro to ggplot)

(Exploratory Data Analysis = EDA)

Exploratory Data Analysis (EDA)

What is EDA and why do we do it?

Before we even do any sort of statistical inference, we need to understand the main characteristics of our dataset.

- ▶ Helps to identify any potential issues or surprising things about our data
- ▶ Helps to check / explore / refine research questions

What is EDA and why do we do it?

EDA is all about asking:


- ▶ What types of variables do we have?
- ▶ Do we have a complete dataset, or do we have missing data or observations?
- ▶ If we have missing data, is it missing equally across observations of different types or concentrated in particular groups?
- ▶ Are there any obvious outliers or strange data points?
- ▶ What do the data 'look' like?
 - ▶ summary measures, measures of centrality, spread
 - ▶ Visualizing the data through plots and tables

Steps of EDA

1. Become familiar with size of data set (number of observations and variables available)
2. What kinds of variables are available
3. For the variables that I'm interested in, are there any missing values or other issues?
4. What does the distribution/frequency of observations look like for the variables I'm interested in? (summary measures, tables and graphs)

Example: TTC subway delays in 2019

- ▶ Data on TTC subway delay times by station and day available from the Open Data Toronto website:
<https://open.toronto.ca/>
- ▶ Accessed through using the opendatatoronto R package:
<https://sharlagelfand.github.io/opendatatoronto/>



Get familiar with dataset

```
delay_2019
```

```
## # A tibble: 19,222 x 11
##   date           time day  station code min_delay min_gap bound line
##   <dtm>          <tim> <chr> <chr>   <chr>      <dbl>   <dbl> <chr> <chr>
## 1 2019-01-01 00:00:00 01:08 Tues~ YORK M~ PUSI         0         0 S    YU
## 2 2019-01-01 00:00:00 02:14 Tues~ ST AND~ PUMST         0         0 <NA> YU
## 3 2019-01-01 00:00:00 02:16 Tues~ JANE    TUSC         0         0 W    BD
## 4 2019-01-01 00:00:00 02:27 Tues~ BLOOR   SUO         0         0 N    YU
## 5 2019-01-01 00:00:00 03:03 Tues~ DUPONT  MUATC        11        16 N    YU
## 6 2019-01-01 00:00:00 03:08 Tues~ EGLINT~ EUATC        11        16 S    YU
## 7 2019-01-01 00:00:00 03:09 Tues~ DUPONT  EUATC         6        11 N    YU
## 8 2019-01-01 00:00:00 03:26 Tues~ ST CLA~ EUATC         4         9 N    YU
## 9 2019-01-01 00:00:00 03:37 Tues~ KENNED~ TUMVS         0         0 E    BD
## 10 2019-01-01 00:00:00 08:04 Tues~ DAVISV~ MUNOA         5        10 S    YU
## # ... with 19,212 more rows, and 2 more variables: vehicle <dbl>,
## #   code_desc <chr>
```


Get familiar with dataset

Dimensions (number of rows x number of columns)

```
dim(delay_2019)
```

```
## [1] 19222    11
```

Variable names

```
colnames(delay_2019)
```

```
## [1] "date"      "time"      "day"       "station"   "code"      "min_delay"  
## [7] "min_gap"   "bound"     "line"      "vehicle"   "code_desc"
```

The summary function is useful for a quick overview

```
summary(delay_2019)
```

```
##      date              time              day
## Min.   :2019-01-01 00:00:00 Length:19222 Length:19222
## 1st Qu.:2019-03-28 00:00:00 Class1:hms   Class :character
## Median :2019-06-27 00:00:00 Class2:difftime Mode :character
## Mean   :2019-06-27 16:58:00 Mode :numeric
## 3rd Qu.:2019-09-25 00:00:00
## Max.   :2019-12-31 00:00:00
##      station      code      min_delay      min_gap
## Length:19222      Length:19222      Min.   : 0.000      Min.   : 0.000
## Class :character      Class :character      1st Qu.: 0.000      1st Qu.: 0.000
## Mode :character      Mode :character      Median : 0.000      Median : 0.000
##                                     Mean   : 2.406      Mean   : 3.536
##                                     3rd Qu.: 3.000      3rd Qu.: 6.000
##                                     Max.   :455.000      Max.   :460.000
##      bound      line      vehicle      code_desc
## Length:19222      Length:19222      Min.   : 0      Length:19222
## Class :character      Class :character      1st Qu.: 0      Class :character
## Mode :character      Mode :character      Median :5239      Mode :character
##                                     Mean   :3974
##                                     3rd Qu.:5671
##                                     Max.   :9206
```

Research question?

- ▶ What are some good potential research questions with this dataset?

Sanity checks

We need to check variables should be what they say they are. If they aren't, the natural next question is to what to do with issues (recode? remove?)

E.g. check days of week make sense with the unique function

```
delay_2019 %>%  
  select(day) %>%  
  unique()
```

```
## # A tibble: 7 x 1  
##   day  
##   <chr>  
## 1 Tuesday  
## 2 Wednesday  
## 3 Thursday  
## 4 Friday  
## 5 Saturday  
## 6 Sunday  
## 7 Monday
```

Sanity checks

Check lines: oh no. some issues here. Some have obvious recodes, others, not so much.

```
delay_2019 %>%  
  select(line) %>%  
  unique() %>%  
  pull() # turn into a vector for better display
```

## [1] "YU"	"BD"	"YU/BD"
## [4] "SHP"	"SRT"	NA
## [7] "YUS"	"B/D"	"BD LINE"
## [10] "999"	"YU/ BD"	"YU & BD"
## [13] "BD/YU"	"YU\\BD"	"46 MARTIN GROVE"
## [16] "RT"	"BLOOR-DANFORTH"	"YU / BD"
## [19] "134 PROGRESS"	"YU - BD"	"985 SHEPPARD EAST EXPR"
## [22] "22 COXWELL"	"100 FLEMINGDON PARK"	"YU LINE"

Data issues

How bad is the mislabeling of lines? look at frequency of cases

NOTE! New very important function: `group_by`

```
delay_2019 %>%  
  group_by(line) %>% # group by line label  
  tally() %>% # count the number of occurrences  
  arrange(-n) # arrange in descending order
```

```
## # A tibble: 24 x 2  
##   line      n  
##   <chr> <int>  
## 1 YU    9275  
## 2 BD    8200  
## 3 SRT    699  
## 4 SHP    600  
## 5 YU/BD   356  
## 6 <NA>    50  
## 7 YU / BD   16  
## 8 YUS      6  
## 9 YU/ BD    3  
## 10 999      2  
## # ... with 14 more rows
```

Missing values

```
delay_2019 %>%  
  summarise_all(.funs = funs(sum(is.na(.))))
```

```
## # A tibble: 1 x 11  
##   date   time   day station   code min_delay min_gap bound  line vehicle  
##   <int> <int> <int>   <int> <int>   <int>   <int> <int> <int>   <int>  
## 1     0     0     0       0     0       0     0 4380    50       0  
## # ... with 1 more variable: code_desc <int>
```

Summary statistics

Most interested in delay minutes, which is the `min_delay` variable

```
delay_2019 %>%  
  summarize(n_obs = n(),  
            mean_delay = mean(min_delay),  
            median_delay = median(min_delay),  
            range_delay = max(min_delay) - min(min_delay),  
            iqr_delay = IQR(min_delay))
```

```
## # A tibble: 1 x 5  
##   n_obs mean_delay median_delay range_delay iqr_delay  
##   <int>   <dbl>         <dbl>     <dbl>     <dbl>  
## 1 18697     2.43           0         455       3
```


Summary statistics

Probably more interesting to do these summaries by line (**stratify** by line); easy extension with the `group_by` function

```
delay_2019 %>%  
  group_by(line) %>%  
  summarize(n_obs = n(),  
            mean_delay = mean(min_delay),  
            median_delay = median(min_delay),  
            range_delay = max(min_delay) - min(min_delay),  
            iqr_delay = IQR(min_delay))
```

```
## # A tibble: 4 x 6  
##   line n_obs mean_delay median_delay range_delay iqr_delay  
##   <chr> <int>     <dbl>         <dbl>         <dbl>     <dbl>  
## 1 BD    8197      2.11            0           180         3  
## 2 SHP    598      2.20            0           165         3  
## 3 SRT    631      5.79            3           284        5.5  
## 4 YU   9271      2.50            0           455         3
```

Summaries

Could also stratify by reason for delay

```
delay_2019 %>%
  group_by(code_desc) %>%
  summarize(n_obs = n(),
            mean_delay = mean(min_delay),
            median_delay = median(min_delay),
            range_delay = max(min_delay) - min(min_delay),
            iqr_delay = IQR(min_delay)) %>%
  arrange(-n_obs)
```

```
## # A tibble: 119 x 6
##   code_desc          n_obs mean_delay median_delay range_delay iqr_delay
##   <chr>          <int>     <dbl>       <dbl>       <dbl>     <dbl>
## 1 Miscellaneous Speed Cont- 1997   0.186         0         19         0
## 2 Injured or ill Customer ~ 1747   0.151         0         54         0
## 3 Operator Overspeeding    1379   0.114         0          8         0
## 4 Passenger Assistance Ala- 1353   0.800         0         12         0
## 5 Disorderly Patron        1147   3.02          3         23         4
## 6 <NA>                931   4.19          0        284         5
## 7 Injured or ill Customer ~ 671   3.92          3         50         5
## 8 Escalator/Elevator Incid- 605   0.00826        0          5         0
## 9 Speed Control Equipment   527   0.436         0         30         0
## 10 ATC Project            514   3.88          3         28         5
## # ... with 109 more rows
```

Summaries

Arrange by mean delay time

```
delay_2019 %>%
  group_by(code_desc) %>%
  summarize(n_obs = n(),
            mean_delay = mean(min_delay),
            median_delay = median(min_delay),
            range_delay = max(min_delay) - min(min_delay),
            iqr_delay = IQR(min_delay)) %>%
  arrange(-mean_delay)
```

```
## # A tibble: 119 x 6
##   code_desc          n_obs mean_delay median_delay range_delay iqr_delay
##   <chr>          <int>     <dbl>     <dbl>     <dbl>     <dbl>
## 1 Traction Power Rail Rela-     1      145         145         0         0
## 2 Priority One - Train in ~    24      78.8         80        193      70.2
## 3 Structure Related Problem     4      70.5         27        228      97.5
## 4 Rail Related Problem         8      58.6          3        455         4
## 5 Fire/Smoke Plan A           6       50        11.5        250      17.5
## 6 Bomb Threat                12      36.7         20        130        32
## 7 Fire/Smoke Plan B - Sour-   84      19.4         11        180      16.2
## 8 Doors Open in Error        11      18.7         16         40         7.5
## 9 Fire/Smoke Plan B - Sour-     2      13.5        13.5         19         9.5
## 10 Suspicious Package        14       13          3.5         67         22
## # ... with 109 more rows
```

EDA: summary so far

- ▶ There's no one checklist of things to look at, depends on your data and research question
- ▶ Get familiar with your dataset
- ▶ Check for missing values, and that existing values make sense
- ▶ Summary statistics depend on your research question of interest
 - ▶ stratifying (`group_by`) by important characteristics often useful

Data visualization

Plot your data!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!

- ▶ We started to compute some summary statistics above, and showed how summaries can be calculated by group and arranged in different ways to get a sense of differences across groups
- ▶ However, graphing/plotting your data is usually the best way to visualize patterns, trends, outliers, issues and other surprising points
- ▶ The most appropriate types of graph for your data depends on:
 - ▶ the type of variable you are interested in (quantitative or qualitative/categorical)
 - ▶ your research questions

Plot your data!!

- ▶ Before you start to do any statistical analysis, you should always plot your data
- ▶ Data visualization is a key part of EDA and essential in understanding the assumptions and outcomes of your eventual statistical analysis

Important types of graphs

Important types of graphs

- ▶ Histograms
- ▶ Bar charts
- ▶ Boxplots
- ▶ Line plots
- ▶ Scatter plots

Example datasets used here


1. TTC subway delays (from above)
2. Country-level indicators, 2009-2017
 - ▶ Uploaded onto Quercus
 - ▶ TFR = total fertility rate
 - ▶ GDP = gross domestics product
 - ▶ dataset also has life expectancy (females), child mortality, maternal mortality

Histograms


Shows the distribution of a **quantitative** variable

- ▶ Histograms show the frequency (count) of observations by value
- ▶ The range of values of a variables is divided into intervals ('bins') and then the number of observations in each bin is tabulated
- ▶ A histogram shows the count of observations in each bin with a rectangle of height equal to the count
- ▶ The x axis is the value bins, the y axis is the count/frequency (or proportion)


Female life expectancy, 2017



Delay times, TTC subway 2019



Making the histogram more informative




Bar charts


Shows summary measures across values of a **categorical** (qualitative) variable

- ▶ Illustrate the value of a particular outcome in a particular category
- ▶ The 'value' can be counts, but could also be a summary measure (e.g. mean)
- ▶ The value is again shown by a rectangle of height equal to the value
- ▶ Bar charts can be plotted vertically or horizontally
- ▶ In the vertical setting, the x axis is the categories and the y axis is the value of the quantitative variable

Number of delays by line, 2019




Same but horizontal




Showing mean delay time

Mean length of delay by line, 2019



More complicated example

Top 5 station by median delay




Box plots

Good for showing summaries of **quantitative** variables across different **categorical** groups.


- ▶ Visualizing quartiles (25/50/75 percentiles) of quantitative data
- ▶ Boxes show the IQR and median
- ▶ Whiskers show values outside IQR (in R/ggplot, default is $1.5 \times \text{IQR}$)
- ▶ Outliers may be shown with individual dots
- ▶ In the vertical case, the x axis is the categories and the y axis is the quantitative variable

Life expectancy (years) by region of the world



Could also do horizontal

Life expectancy (years) by region of the world




Line plots

Best used to describe values of a **quantitative** variable (on y axis) across sequential values of another **quantitative** variable on the x axis

- ▶ Plots a series of values of a quantitative variable connected together by a line
- ▶ Useful to visualize trends over time

Total fertility rate, Canada and the US

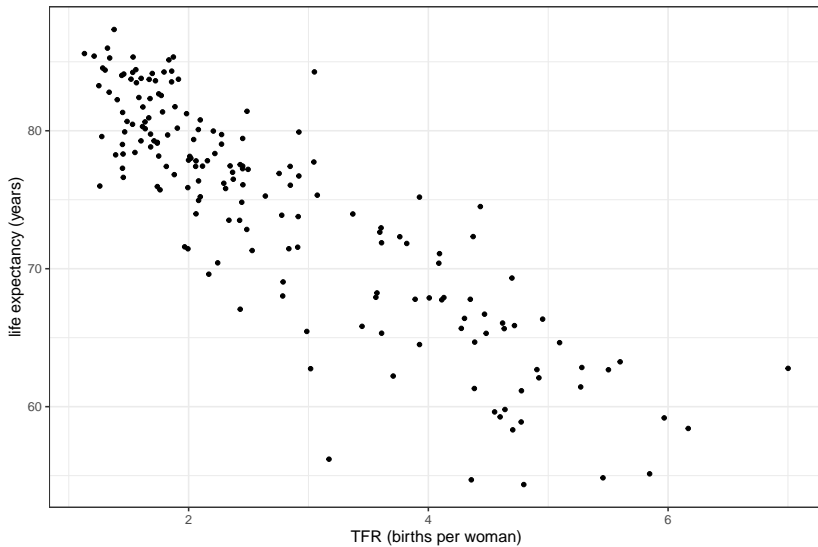


Scatter plots


Shows relationship between two different **quantitative** variables

- ▶ Uses dots to represent values for two different **quantitative** values
- ▶ The position of each dot on the x and y axis indicates values for an individual data point
- ▶ Extremely useful in visualizing the relationship between two quantitative variables

TFR versus life expectancy, 2017



TFR versus life expectancy, 2017



Aside: another summary measure

Based on the previous graphs, evidence to suggest a relationship between TFR and life expectancy

- ▶ as TFR goes up, life expectancy tends to go down
- ▶ life expectancy is **negatively correlated** with TFR

Correlation is the statistical measure of the relationship between two variables. **Pearson's correlation coefficient**, r_{xy} summarizes this relationship into one number. For an observation sample of two random variables x_1, x_2, \dots, x_n and y_1, y_2, \dots, y_n ,

$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

Correlation coefficients in R


Easiest to use the function `cor`

```
country_ind_2017 <- country_ind %>% filter(year==2017)
country_ind_2017 %>%
  select(tfr, life_expectancy) %>%
  summarize(correlation = cor(tfr, life_expectancy))
```

```
## # A tibble: 1 x 1
##   correlation
##         <dbl>
## 1      -0.868
```

```
# alternative code
# cor(country_ind_2017$tfr, country_ind_2017$life_expectancy)
```


TFR versus life expectancy, 2017



Plot your data!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!













The correlation coefficient tells us in a single number that there is a negative relationship observed between TFR and life expectancy. So why bother plotting at all?

Anscombe's quartet







X Mean: 54.26
 Y Mean: 47.83
 X SD : 16.76
 Y SD : 26.93
 Corr. : -0.06



Source: <https://www.autodeskresearch.com/publications/samestats>

Where are the pie charts?

Don't use pie charts!, Humans are inherently bad at judging angles, which is what you have to do with a pie chart. Use a bar chart instead.



Introduction to ggplot


ggplot

- ▶ ggplot is the graphing package that goes with the tidyverse in R
- ▶ Very powerful to make a wide range of graphics
- ▶ Every graph so far this lecture was done in ggplot
- ▶ ggplot code works in layers, with each layer adding complexity
 - ▶ start with defining dataset and different variables
 - ▶ add on type of plot
 - ▶ scales
 - ▶ layout (facets)
 - ▶ themes, fonts, sizes...

More practice in lab, but here's a starting example

Reproducing the TFR versus life expectancy chart, colored by region

TFR versus life expectancy, 2017



Data

```
# read in the data
```

```
country_ind <- read_csv("../data/country_indicators.csv")  
country_ind
```

```
## # A tibble: 1,584 x 9
```

```
##   country_code country region year  tfr life_expectancy child_mort  
##   <chr>         <chr>   <chr> <dbl> <dbl>         <dbl>    <dbl>  
## 1 AFG         Afghan~ South~ 2009  6.18         61.9      93.9  
## 2 AFG         Afghan~ South~ 2010  5.98         62.5      90.0  
## 3 AFG         Afghan~ South~ 2011  5.77         63        86.3  
## 4 AFG         Afghan~ South~ 2012  5.56         63.5      82.9  
## 5 AFG         Afghan~ South~ 2013  5.36         64.0      79.6  
## 6 AFG         Afghan~ South~ 2014  5.16         64.5      76.6  
## 7 AFG         Afghan~ South~ 2015  4.98         64.9      73.8  
## 8 AFG         Afghan~ South~ 2016  4.80         65.3      71.2  
## 9 AFG         Afghan~ South~ 2017  4.63         65.7      68.8  
## 10 ALB        Albania Devel~ 2009  1.65         79.0      16.7  
## # ... with 1,574 more rows, and 2 more variables: maternal_mort <dbl>,  
## #   gdp <dbl>
```

```
# filter to just be 2017
```


```
country_ind_2017 <- country_ind %>% filter(year==2017)
```

A blank canvas

`aes` stands for aesthetic and tells `ggplot` the main characteristics of your plot (`x`, `y`, and if the color or fill vary by group)

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy))
```

```
#print  
plot1
```




Add the points

Add layers with ggplot using the +

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy)) +  
  geom_point()
```


plot1



Tidy up labels

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy)) +  
  geom_point() +  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)")
```


plot1



Title

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy)) +  
  geom_point()+  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)") +  
  ggtitle("TFR versus life expectancy, 2017")
```

plot1




Change color of points

to see all colors, type `colors()`

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy)) +  
  geom_point(color = "red")+  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)") +  
  ggtitle("TFR versus life expectancy, 2017")
```

plot1




Coloring by group

This goes in the `aes()` because it **depends on the data**

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy, color = region)) +  
  geom_point() +  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)") +  
  ggtitle("TFR versus life expectancy, 2017")  
  
plot1
```

TFR versus life expectancy, 2017




Change theme (optional) and size of points

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy, color = region)) +  
  geom_point(size = 2) +  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)") +  
  ggtitle("TFR versus life expectancy, 2017") +  
  theme_bw(base_size = 14)
```

plot1

TFR versus life expectancy, 2017




Change color scheme

viridis and brewer both good options

```
plot1 <- ggplot(data = country_ind_2017, aes(x = tfr, y = life_expectancy, color = region)) +  
  geom_point(size = 2) +  
  xlab("TFR (births per woman)") +  
  ylab("life expectancy (years)") +  
  ggtitle("TFR versus life expectancy, 2017") +  
  theme_bw(base_size = 14) +  
  scale_color_viridis_d()
```

plot1

TFR versus life expectancy, 2017



Summary

- ▶ EDA and data visualization is often just as informative and important as statistical analysis
- ▶ It is essential to understand the structure of your data, missing-ness, any outliers/issues, and the raw patterns in your data before deciding on your statistical analysis
- ▶ Plot, plot, plot
- ▶ Practice, practice, practice

Summary

Plots:

- ▶ Bar charts for categorical/qualitative variables
- ▶ Histograms, boxplots for one quantitative variable (potentially across multiple categories)
- ▶ Line plots and scatter plots for two quantitative variables (line plot when one is sequential)

Lab

- ▶ Quick overview of the end of last week's lab (see video)
- ▶ Practice with `ggplot` and how to graph important types of plots in R.

Data ideas

- ▶ IPUMS: <https://ipums.org/>
- ▶ ICPSR:
<https://www.icpsr.umich.edu/web/pages/ICPSR/thematic-collections.html>
- ▶ CHASS SDA: <https://datacentre.chass.utoronto.ca/>
- ▶ Toronto Open Data Portal: <https://open.toronto.ca/> or use opendatatoronto R package (ask for code)
- ▶ UN WPP: <https://population.un.org/wpp/>
- ▶ NBER:
<https://www.nber.org/research/data?page=1&perPage=50>